



**SARDIGNA CHIRCAS
SARDEGNA RICERCHÉ**

open:campus

Artificial Intelligence for developers

8 weekend per diventare Machine Learning Specialist



Natural Language Processing

Maurizio Atzori

Università degli Studi di Cagliari

atzori@unica.it

February 9-10 , 2024



Temptative (outline) of the course

- **Intro on AI, ML and NLP**
- **Text Processing**
- **Words and Corpora**
- **Lexical similarity**
- **Language Modeling**
- **Text Classification**
- **Semantic similarity**
- **Knowledge Graphs**
- **Intro to Large Language Models**

An abstract background on the left side of the slide. It features a dark blue/black field with glowing yellow and white lines forming a network or web. In the center of this network, the letters 'IA' are prominently displayed in a stylized, wireframe font. The background also contains faint, scattered binary code (0s and 1s) and some blurred text fragments.

Regular Expressions

Text Processing

Regular expressions

A formal language for specifying text strings

How can we search for any of these?

- woodchuck
- woodchucks
- Woodchuck
- Woodchucks



Regular Expressions: Disjunctions

Letters inside square brackets []

Pattern	Matches
[wW]oodchuck	Woodchuck, woodchuck
[1234567890]	Any digit

Ranges [A-Z]

Pattern	Matches	
[A-Z]	An upper case letter	<u>D</u> renched Blossoms
[a-z]	A lower case letter	<u>m</u> y beans were impatient
[0-9]	A single digit	Chapter <u>1</u> : Down the Rabbit Hole

Regular Expressions: Negation in Disjunction

Negations [^Ss]

- Carat means negation only when first in []

Pattern	Matches	
[^A-Z]	Not an upper case letter	O <u>y</u> fn pripetchik
[^Ss]	Neither 'S' nor 's'	<u>I</u> have no exquisite reason"
[^e^]	Neither e nor ^	Look <u>h</u> ere
a^b	The pattern a carat b	Look up <u>a^b</u> now

Regular Expressions: More Disjunction

Woodchuck is another name for groundhog!

The pipe | for disjunction

Pattern	Matches
<code>groundhog woodchuck</code>	woodchuck
<code>yours mine</code>	yours
<code>a b c</code>	= <code>[abc]</code>
<code>[gG]roundhog [Ww]oodchuck</code>	Woodchuck



Regular Expressions: ? *+ .

Pattern	Matches	
<code>colou?r</code>	Optional previous char	<u>color</u> <u>colour</u>
<code>oo*h!</code>	0 or more of previous char	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
<code>o+h!</code>	1 or more of previous char	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
<code>baa+</code>		<u>baa</u> <u>baaa</u> <u>baaaa</u> <u>baaaaa</u>
<code>beg.n</code>		<u>begin</u> <u>begun</u> <u>begun</u> <u>beg3n</u>



Stephen C
Kleene

Kleene *, Kleene +

Regular Expressions: Anchors ^ \$

Pattern	Matches
<code>^[A-Z]</code>	<u>P</u> alo Alto
<code>^[^A-Za-z]</code>	<u>1</u> "Hello"
<code>\. \$</code>	The end <u>.</u>
<code>. \$</code>	The end <u>?</u> The end <u>!</u>

Example

Find me all instances of the word “the” in a text.

the

Misses capitalized examples

[tT]he

Incorrectly returns other or theology

[^a-zA-Z][tT]he[^a-zA-Z]

Errors

The process we just went through was based on fixing two kinds of errors:

1. Matching strings that we should not have matched
(there, then, other)

False positives (Type I errors)

2. Not matching things that we should have matched
(The)

False negatives (Type II errors)

Errors cont.

In NLP we are always dealing with these kinds of errors.

Reducing the error rate for an application often involves two antagonistic efforts:

- **Increasing accuracy or precision** (minimizing false positives)
- **Increasing coverage or recall** (minimizing false negatives).

Summary

Regular expressions play a surprisingly large role

- Sophisticated sequences of regular expressions are often the first model for any text processing text

For hard tasks, we use machine learning classifiers

- But regular expressions are still used for pre-processing, or as features in the classifiers
- Can be very useful in capturing generalizations

An abstract background on the left side of the slide. It features a dark blue/black field with glowing yellow and white lines forming a network or web. In the center of this network, the letters 'IA' are prominently displayed in a stylized, wireframe font. The background also contains faint, scattered binary code (0s and 1s) and some blurred text fragments.

Regular Expressions

Text Processing

An abstract graphic on the left side of the slide. It features a dark background with glowing binary code (0s and 1s) and a network of white lines connecting various points, resembling a data or neural network. The graphic is partially obscured by an orange triangle at the top and a dark blue triangle at the bottom.

Substitutions and ELIZA

Text Processing

Substitutions

Substitution in Python and UNIX commands:

```
s/regexpl/pattern/
```

e.g.:

```
s/colour/color/
```

Capture Groups

- Say we want to put angles around all numbers:

the 35 boxes → *the <35> boxes*

- Use parens () to "capture" a pattern into a numbered register (1, 2, 3...)
 - Use \1 to refer to the contents of the register
- `s / ([0-9] +) / <\1> /`

Capture groups: multiple registers

`/the (.*)er they (.*) , the \1er we \2/`

Matches

the faster they ran, the faster we ran

the better they behave, the better we behave

But not

the faster they ran, the faster we behave

But suppose we don't want to capture?

Parentheses have a double function: grouping terms, and capturing

Non-capturing groups: add a ?: after paren:

`/(?:some|a few) (people|cats) like some \1/`
matches

- some cats like some cats

but not

- some cats like some some

Lookahead assertions

`(?= pattern)` is true if pattern matches, but is **zero-width; doesn't advance character pointer**

`(?! pattern)` true if a pattern does not match

How to match, at the beginning of a line, any single word that doesn't start with "Volcano":

```
/^(?!Volcano) [A-Za-z]+/
```

Simple Application: ELIZA

Early NLP system that imitated a Rogerian psychotherapist

- Joseph Weizenbaum, 1966.

Uses pattern matching to match, e.g.,:

- "I need X"

and translates them into, e.g.

- "What would it mean to you if you got X?"

Simple Application: ELIZA

Men are all alike.

IN WHAT WAY

They're always bugging us about something or other. CAN YOU THINK OF A SPECIFIC

EXAMPLE

Well, my boyfriend made me come here.

YOUR BOYFRIEND MADE YOU COME HERE

He says I'm depressed much of the time.

I AM SORRY TO HEAR YOU ARE DEPRESSED

How ELIZA works

s/. * I'M (depressed|sad) . */ I AM SORRY TO HEAR YOU ARE \1/

s/. * I AM (depressed|sad) . */ WHY DO YOU THINK YOU ARE \1/

s/. * all . */ IN WHAT WAY?/

s/. * always . */ CAN YOU THINK OF A SPECIFIC EXAMPLE?/

An abstract background on the left side of the slide. It features a dark blue/black field with glowing yellow and white lines forming a complex network or graph. In the center of this network, the letters 'IA' are prominently displayed in a stylized, wireframe font. The background also contains faint, scattered binary code (0s and 1s) and some blurred text fragments.

Substitutions and ELIZA

Text Processing

Exercise on Regex